

Ansible Assignment

Cloud Infrastructure

Rob Shelly
20068406

April 11, 2018


Contents

1	Introduction	3
2	Ansible	3
3	Roles	4
3.1	Provision EC2 Instance	4
3.2	Nginx	5
3.3	HAProxy	6
4	Playbooks	7
5	Ansible Vault	8
6	Running the Playbook	8
A	Project Repository	10
B	Provision EC2 Instance Variable files	11
C	Provision EC2 Instance Variable files	12

1 Introduction

The document will outline the concepts behind and the tasks required to configure and web server on AWS using Ansible. The web server will consist of two Nginx servers running in a private subnet, with a HAProxy load balancer running in a public subnet balancing traffic to the web servers. Also running in the public subnet is the Ansible server which will provision the Nginx and HAProxy servers. This is illustrated in [Figure 1](#). For the purposes of this exercise, the Nginx server have been configured to redirect to *serversforhackers.com*.

Figure 1: System Architecture



2 Ansible

Ansible configures servers through the use of a number of key concepts, the first of which is *Playbooks*. Playbooks define the configuration to be deployed to a server. When executed against a server they run a number of tasks which will make the necessary installations, create the necessary file etc. to configure the server. They are also idempotent, meaning that they can be run multiple times without changing or breaking the initial result.

Playbooks use *tasks* and *handlers* to configure servers. Tasks are the individual steps required for configuration, such as adding a package repository or installing a package. When a playbook is run, all of its tasks are executed in order. Handlers, however, are only executed when called. For example, a handler to start a service might be called after installing the service.

Writing the entire configuration for a server in one playbook could be quite complex and would not be very reusable. This is where *roles* are used. Roles group tasks, handlers and other files such as variable files and templates

together in a standard directory structure. A role can then be called by a playbook. A role can even be called but instructed to use a different variable file. This allows for greater reuse in Ansible.

3 Roles

In order to deploy the system described above, Ansible must create EC2 instances within the VPC, then provision the necessary servers with Nginx or HAProxy. Therefore, three roles will be used to create this system:

- Provision EC2 Instance
- Nginx
- HAProxy

3.1 Provision EC2 Instance

This role is used to provision instances within EC2. The role defines three main tasks:

- Provision the instance(s)
- Add the instances to the dynamic list of hosts held by Ansible
- Wait for the instance(s) to be reachable via SSH

These tasks are shown below [Code Sample 1](#). There are a number of aspects to note in these tasks.

- There are a large number of variables included in the tasks. This will allow the role to be reused. For example, the *instance_count* will allow the role to be used to create two Nginx servers initially but only one HAProxy server later.
- The task adds the instance(s) it created to a group in the hosts file, defined by the *host_group* variable. This means that instances created by the role can then be provisioned with the correct configuration by running playbooks against the hosts in the *host_group*.
- The task adds the private IP addresses of the instance to the hosts file. This is necessary because some of the servers are created in a private subnet and therefore have no public IP address. However, as the Ansible server is running in the public subnet, it can connect to any server in either the private or public subnets using their private IP addresses.

In order to use the role to provision servers differently, (i.e. for Nginx server or a HAProxy server) two different variable files are defined; *webserver.yml* and *loadbalancer.yml*. These are shown in [appendix B](#).

Code Sample 1: provision-ec2-instance/tasks/main.yml

```
---
- name: Provision EC2 Box
  local_action:
  module: ec2
  profile: "{{ aws_profile }}"
  key_name: "{{ ec2_keypair }}"
  group_id: "{{ ec2_security_group }}"
  instance_type: "{{ ec2_instance_type }}"
  image: "{{ ec2_image }}"
  vpc_subnet_id: "{{ ec2_subnet_ids|random }}"
  region: "{{ ec2_region }}"
  instance_tags: '{"Name": "{{ ec2_tag_Name }}", "Type": "{{ ec2_tag_Type }}", "Environment": "{{ ec2_tag_Environment }}"}'
  assign_public_ip: "{{ assign_public_ip }}"
  wait: true
  count: "{{ instance_count }}"
  volumes:
  - device_name: /dev/sda1
    device_type: gp2
    volume_size: "{{ ec2_volume_size }}"
    delete_on_termination: true
  register: ec2

- debug: var=item
  with_items: ec2.instances

- name: Add instances to host group
  add_host:
  hostname: "{{ item.private_ip }}"
  ansible_user: "{{ ssh_user }}"
  ansible_ssh_private_key_file: "{{ ssh_key_path }}"
  ansible_python_interpreter: /usr/bin/python3
  ansible_ssh_extra_args: '-o StrictHostKeyChecking=no'
  groups: "{{ host_group }}"
  with_items: "{{ ec2.instances }}"

- name: Wait for the instances to boot by checking the ssh port
  wait_for: host={{ item.private_ip }} port=22 delay=60 timeout=320 state=started
  with_items: "{{ ec2.instances }}"
```

3.2 Nginx

The Nginx role consists of the following main elements:

- **Tasks:** These are the main tasks required to configure Nginx. They include adding the package repository

and installing Nginx, and moving the necessary configuration files into place.

- **Handlers:** The Nginx role defines two handlers: *Start Nginx* and *Reload Nginx*. The are called as required by the main tasks. For example, after creating the Nginx configuration, the *Reload Nginx* handler is called.
- **Templates:** Templates can be used to defined files that must be copied to the server. However they also allow for variable through the use of Jinja. The Nginx roles use a template to create the Nginx configuration file. This can be seen in [Code Sample 2](#).
- **Variables:** The Nginx role also defines a variables files. In this case the web server domain is defined as a variable.

Code Sample 2: nginx/templates/serverforhacker.com.conf.j2

```
server {
    # Enforce the use of HTTPS
    listen 80 default_server;
    server_name {{ domain }};
    return 301 https://$server_name$request_uri;
}
server {
    listen 443 ssl default_server;
    root /var/www/{{ domain }}/public/;
    index index.html index.htm index.php;
    access_log /var/log/nginx/{{ domain }}.log;
    error_log /var/log/nginx/{{ domain }}-error.log error;
    server_name {{ domain }};
    charset utf-8;
    include h5bp/basic.conf;
    include h5bp/directive-only/ssl.conf;

    location / {
        try_files $uri $uri/ /index.php$is_args$args;
    }
    location = /favicon.ico { log_not_found off; access_log off; }
    location = /robots.txt { log_not_found off; access_log off; }
}
```

3.3 HAProxy

The HAProxy roles is similar to the Nginx role. It defines two main tasks: *Install HAProxy* and *Update HAProxy Config*. Like, Nginx it also defines handlers to start and reload HAProxy. The HAProxy configuration is provided by a template file. This is shown below in [Code Sample 3](#).

Of particular is the *server* block. This is adding the hosts in the *webservers* group and their IP addresses to the HAProxy server configuration. As stated above, when the Provision EC2 Instance role runs, it adds the instance to the Ansible hosts. When run with the *webservers* variables file, (which is used to provision the Nginx servers), the instances are added under the *webservers* group. Therefore, this section of the template is configuring any provisioned Nginx instances as the server for HAProxy.

Code Sample 3: haproxy/templates/haproxy.cfg

```
global
    log 127.0.0.1 local0 notice
    maxconn 2000
    user haproxy
    group haproxy

defaults
    log global
    mode http
    option httplog
    option dontlognull
    retries 3
    option redispatch
    timeout connect 5000
    timeout client 10000
    timeout server 10000

listen {{haproxy_app_name}}
    bind 0.0.0.0:80
    mode {{haproxy_mode}}
    stats {{haproxy_enable_stats}}
    {% if haproxy_enable_stats == 'enable' %}
    stats uri /haproxy_stats
    stats realm Strictly\ Private
    {% for user in haproxy_stats_users %}
    stats auth {{user.username}}:{{user.password}}
    {% endfor %}
    {% endif %}
    balance {{haproxy_algorithm}}
    option httpclose
    option forwardfor

    {% for host in groups['webservers'] %}
    server {{hostvars[host]['ansible_hostname']}} {{ hostvars[host]['ansible_eth0']['ipv4
        ']['address'] }}:80 check
    {% endfor %}
```

4 Playbooks

With all the roles defined, playbooks can be created to run these roles against the necessary servers. Although one playbook could be created to run each of the roles, it is preferable to create multiple playbooks for each of the server to be configured and include these in one *master* playbook. These playbooks are included in appendix C. The playbooks are then called from one master playbook, shown below in [Code Sample 4](#)

Code Sample 4: main-playbook.yml

```
---
- import_playbook: provision-ec2.yml type='webserver'
- import_playbook: nginx.yml hosts='webserver'
- import_playbook: provision-ec2.yml type='loadbalancer'
- import_playbook: haproxy.yml hosts='loadbalancer'
```

Here, the Provision EC2 playbook is called twice. However, in order to deploy the different instances (i.e. two in the private subnet for Nginx, one in the public subnet for HAProxy), the *type* variable is used to specify which variables files should be used. Also, the *hosts* variables are used to instruct the Nginx and HAProxy playbooks which server they should run against.

5 Ansible Vault

Each of the roles have specified variable files. These have been used to allow for greater reuse of the roles. However, they can also contain sensitive information. Ansible Vault provides a method of encrypting these files while still allowing them to be used by playbooks, by providing a password. All of the variable files were encrypted using the following command:

Code Sample 5: vault.yml

```
ansible-vault encrypt roles/<rolename>/vars/<filename>
```

They can also be edited or decrypted using the *edit* and *decrypt* flags.

6 Running the Playbook

The *main-playbook* can now be run. However, Ansible must have SSH access to the server it created in the private subnet in order to install Nginx on them. Therefore, the playbook was run on the Ansible server within the VPC, specifically within the public subnet. The following command runs the playbook, asking the user to enter their password to use the variable files encrypted by Ansible vault.

Code Sample 6: run.yml

```
ansible-playbook --ask-vault-pass -i localhost, main-playbook.yml
```

Once the playbook visiting the AWS console verified that the correct instances have been created, shown in ???. Note that the Nginx server have no public IP addresses as they are within the private subnet.

Figure 2: EC2 Instance

<input type="checkbox"/>	Name	Instance Type	Instance State	IPv4 Public IP	Security Groups	Subnet ID
<input checked="" type="checkbox"/>	ansible-server	t2.micro	● running	34.244.160.15	load-balancer-sg	subnet-a412b9ec
<input type="checkbox"/>	ansible-lab-webserver	t2.micro	● running	-	web-servers-sg	subnet-6f1cb727
<input type="checkbox"/>	ansible-lab-webserver	t2.micro	● running	-	web-servers-sg	subnet-6f1cb727
<input type="checkbox"/>	ansible-lab-loadbalancer	t2.micro	● running	52.30.254.37	load-balancer-sg	subnet-a412b9ec

By visiting the *stats* page of the HAProxy server the status of the two web servers can be verified.

Figure 3: HAProxy

ansible-lab												
	Queue			Session rate			Sessions					
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last
Frontend				1	4	-	1	4	2 000	10		
ip-10-0-1-158	0	0	-	0	1		0	1	-	1	1	1m38s
ip-10-0-1-173	0	0	-	0	1		0	1	-	1	1	35s
Backend	0	0		0	1		0	1	200	2	2	0s

Appendix A Project Repository

Relevant code samples have been provided in this document. However the project in it's entirety can be found at the following link:

<https://github.com/robshelly/ansible-lab>

Appendix B Provision EC2 Instance Variable files

Code Sample 7: provision-ec2-instance/vars/webserver.yml

```
---
ec2_keypair: "ansible"
ec2_security_group: "sg-73dbc209"
ec2_instance_type: "t2.micro"
ec2_image: "ami-f90a4880"
ec2_subnet_ids: ['subnet-6f1cb727']
ec2_region: "eu-west-1"
ec2_tag_Name: "ansible-lab-webserver"
ec2_tag_Type: "ansible-lab-webserver"
ec2_tag_Environment: "production"
ec2_volume_size: 8
assign_public_ip: no
host_group: webserver
instance_count: 2
ssh_user: ubuntu
ssh_key_path: ~/keypairs/ansible.pem
aws_profile: ansible
```

Code Sample 8: provision-ec2-instance/vars/loadbalance.yml

```
---
ec2_keypair: "ansible"
ec2_security_group: "sg-6cdcc516"
ec2_instance_type: "t2.micro"
ec2_image: "ami-f90a4880"
ec2_subnet_ids: ['subnet-a412b9ec']
ec2_region: "eu-west-1"
ec2_tag_Name: "ansible-lab-loadbalancer"
ec2_tag_Type: "ansible-lab-loadbalancer"
ec2_tag_Environment: "production"
ec2_volume_size: 8
assign_public_ip: yes
host_group: loadbalancer
instance_count: 1
ssh_user: ubuntu
ssh_key_path: ~/keypairs/ansible.pem
aws_profile: ansible
```

Appendix C Provision EC2 Instance Variable files

Code Sample 9: provision-ec2.yml

```
---
- hosts: localhost
  connection: local
  gather_facts: false
  user: root
  pre_tasks:
  - include_vars: ~/ansible-lab/roles/provision-ec2-instance/vars/{{type}}.yaml
  roles:
  - provision-ec2-instance
```

Code Sample 10: nginx.yml

```
---
- hosts: "{{hosts}}"
  become: yes
  become_user: root
  roles:
  - nginx
```

Code Sample 11: haproxy.yml

```
---
- hosts: "{{hosts}}"
  become: yes
  become_user: root
  roles:
  - haproxy
```